

L-Eval | GPU staging occupancy | concurrency 8 | all 3 repeats

Qwen3-14B | 72 requests/run | 2 DRAM prefetch workers | 18 documents x 4 questions

Repeat 1 Repeat 2 Repeat 3 Solid: sampled peak per 2s bin | Dashed: sampled mean per 2s bin
Y: allocated staging / configured capacity (%) X: seconds since each run started



All panels share the same axes. Runs end at different times; no zero padding or time normalization.
Sampled peaks may miss brief allocation spikes. Occupancy is not GPU compute utilization.
OFF disables only DRAM prefetch; DAOS prefetch remains enabled. Equal times do not imply equal requests.

L-Eval | GPU staging occupancy | concurrency 16 | all 3 repeats

Qwen3-14B | 72 requests/run | 2 DRAM prefetch workers | 18 documents x 4 questions

Repeat 1 Repeat 2 Repeat 3 Solid: sampled peak per 2s bin | Dashed: sampled mean per 2s bin
Y: allocated staging / configured capacity (%) X: seconds since each run started



All panels share the same axes. Runs end at different times; no zero padding or time normalization.
Sampled peaks may miss brief allocation spikes. Occupancy is not GPU compute utilization.
OFF disables only DRAM prefetch; DAOS prefetch remains enabled. Equal times do not imply equal requests.

L-Eval | GPU staging occupancy | concurrency 8 | repeat 1

Qwen3-14B | 72 requests/run | 2 DRAM prefetch workers | 18 documents x 4 questions

Repeat 1 Repeat 2 Repeat 3 Solid: sampled peak per 2s bin | Dashed: sampled mean per 2s bin
Y: allocated staging / configured capacity (%) X: seconds since each run started

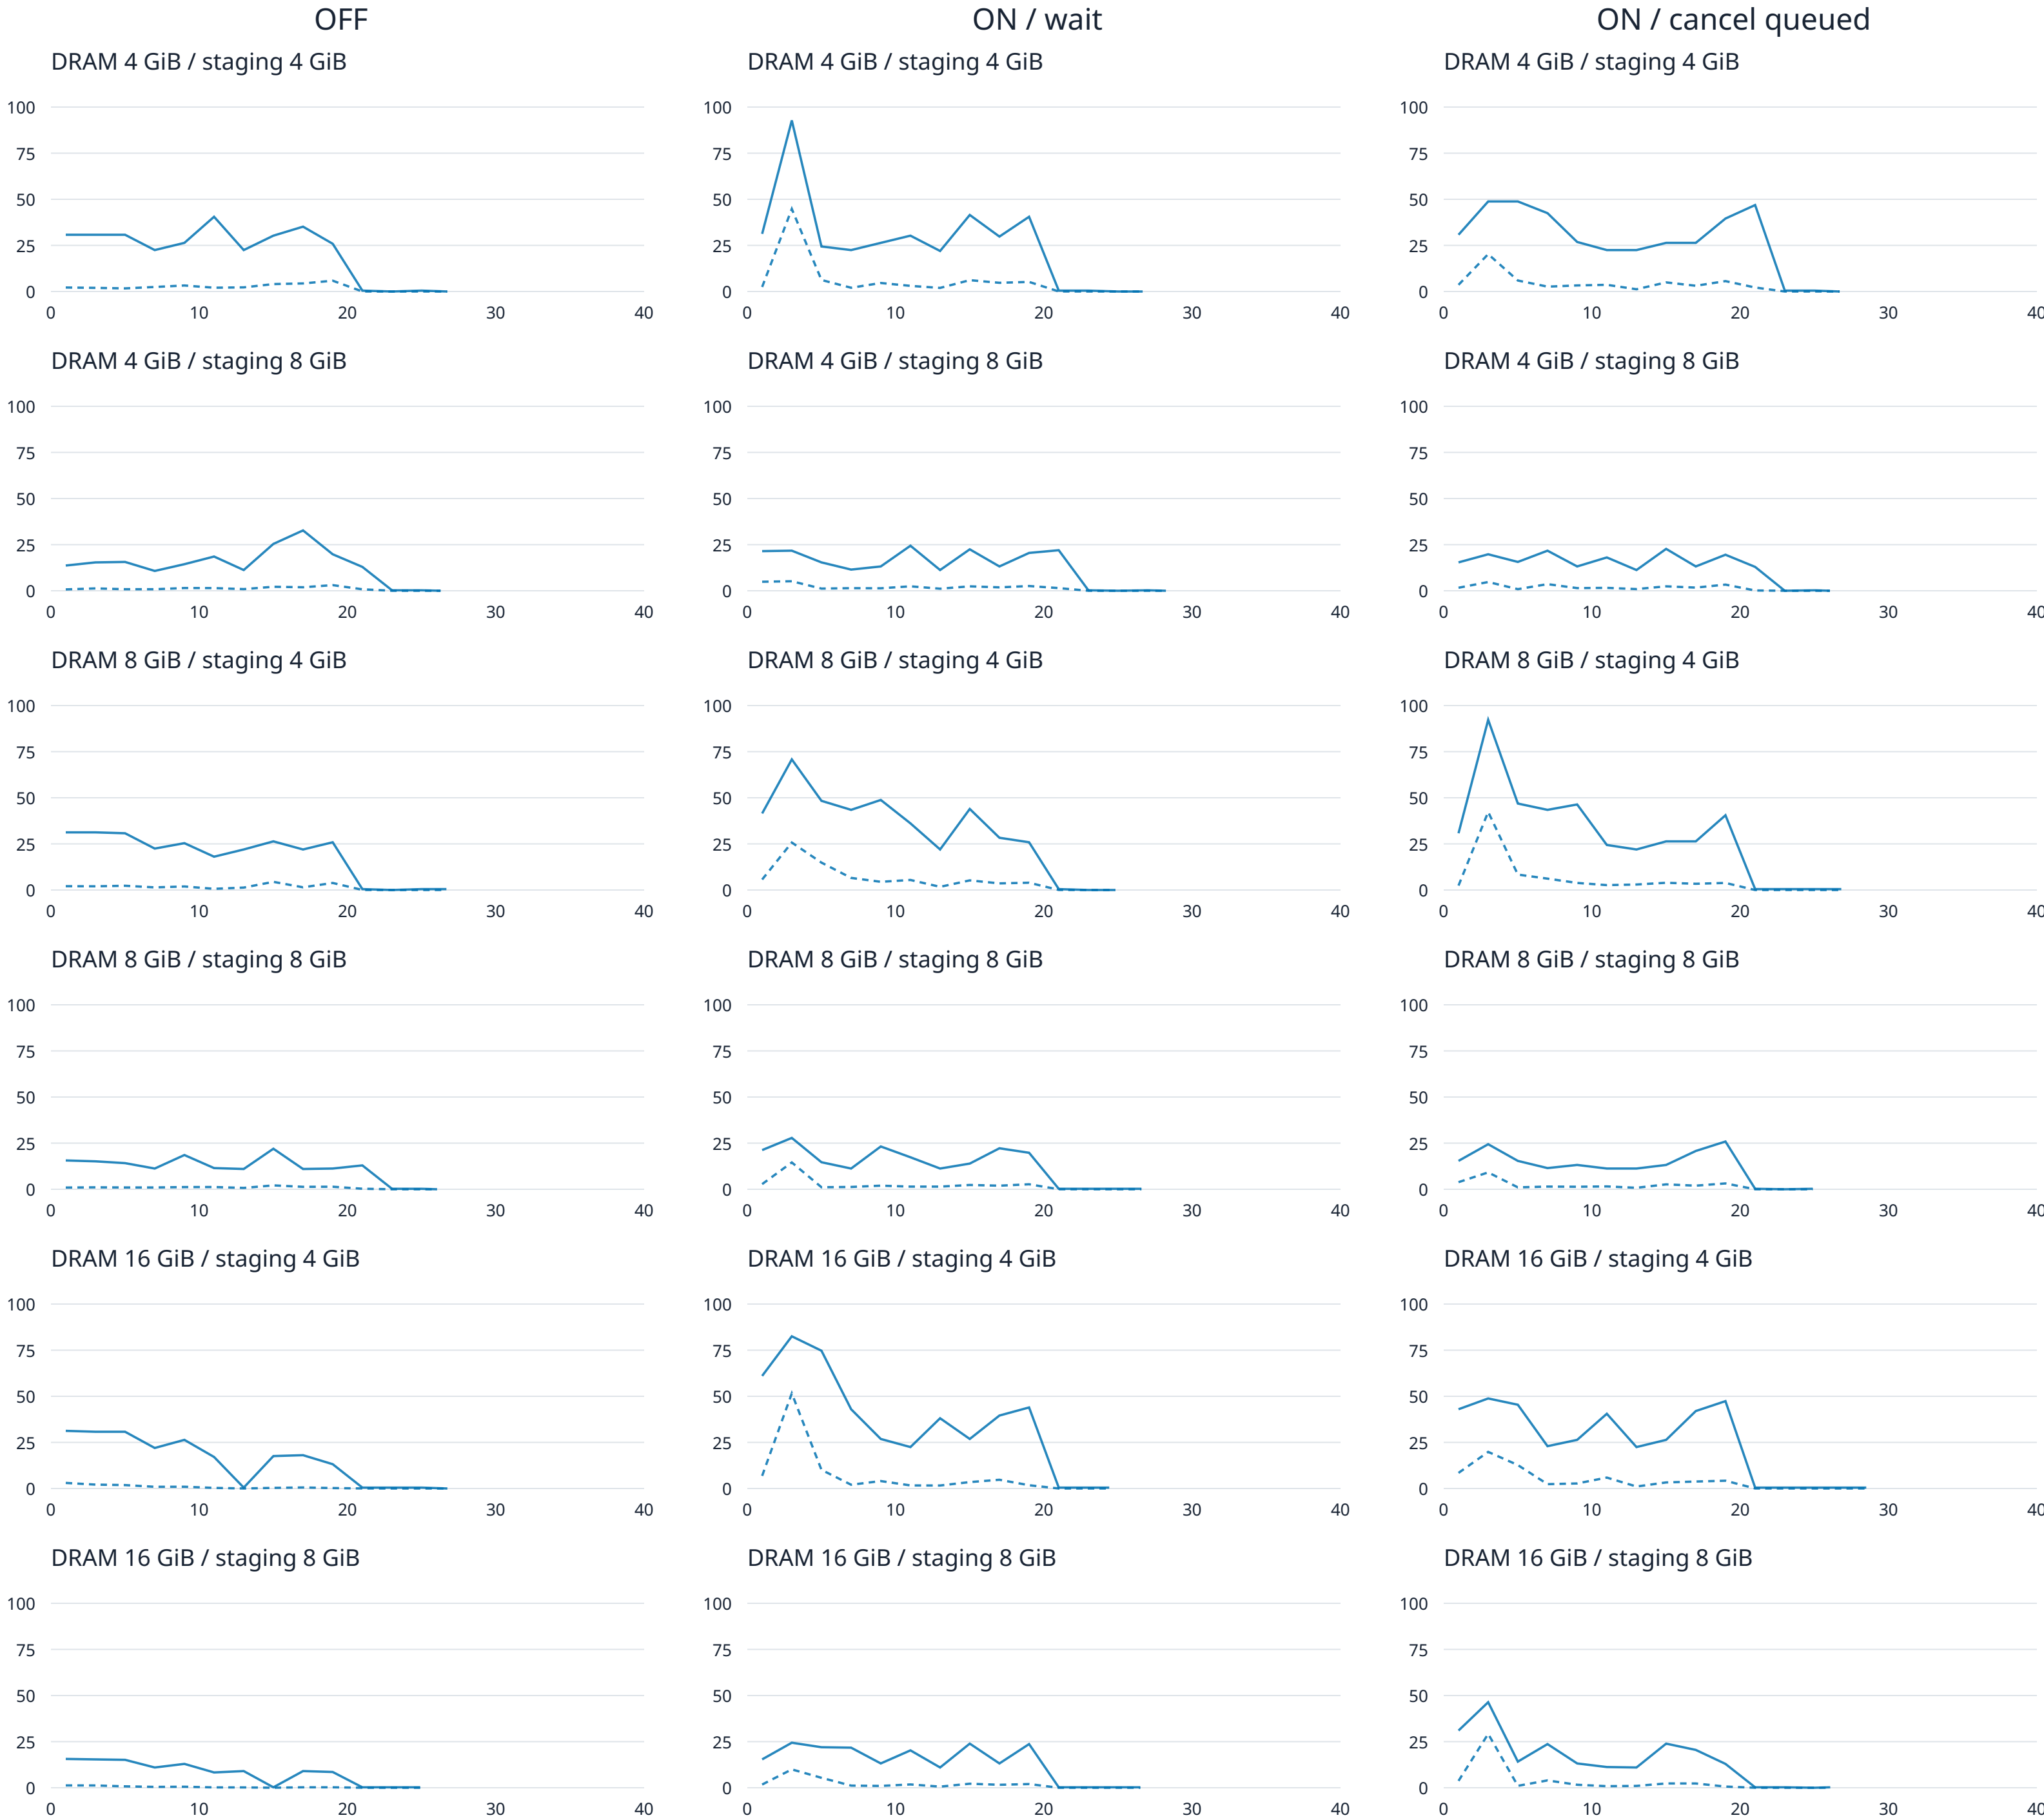


All panels share the same axes. Runs end at different times; no zero padding or time normalization.
Sampled peaks may miss brief allocation spikes. Occupancy is not GPU compute utilization.
OFF disables only DRAM prefetch; DAOS prefetch remains enabled. Equal times do not imply equal requests.

L-Eval | GPU staging occupancy | concurrency 16 | repeat 1

Qwen3-14B | 72 requests/run | 2 DRAM prefetch workers | 18 documents x 4 questions

Repeat 1 Repeat 2 Repeat 3 Solid: sampled peak per 2s bin | Dashed: sampled mean per 2s bin
Y: allocated staging / configured capacity (%) X: seconds since each run started



All panels share the same axes. Runs end at different times; no zero padding or time normalization.
Sampled peaks may miss brief allocation spikes. Occupancy is not GPU compute utilization.
OFF disables only DRAM prefetch; DAOS prefetch remains enabled. Equal times do not imply equal requests.

L-Eval | GPU staging occupancy | concurrency 8 | repeat 2

Qwen3-14B | 72 requests/run | 2 DRAM prefetch workers | 18 documents x 4 questions

Repeat 1 Repeat 2 Repeat 3 Solid: sampled peak per 2s bin | Dashed: sampled mean per 2s bin
Y: allocated staging / configured capacity (%) X: seconds since each run started



All panels share the same axes. Runs end at different times; no zero padding or time normalization.
Sampled peaks may miss brief allocation spikes. Occupancy is not GPU compute utilization.
OFF disables only DRAM prefetch; DAOS prefetch remains enabled. Equal times do not imply equal requests.

L-Eval | GPU staging occupancy | concurrency 16 | repeat 2

Qwen3-14B | 72 requests/run | 2 DRAM prefetch workers | 18 documents x 4 questions

Repeat 1 Repeat 2 Repeat 3 Solid: sampled peak per 2s bin | Dashed: sampled mean per 2s bin
Y: allocated staging / configured capacity (%) X: seconds since each run started

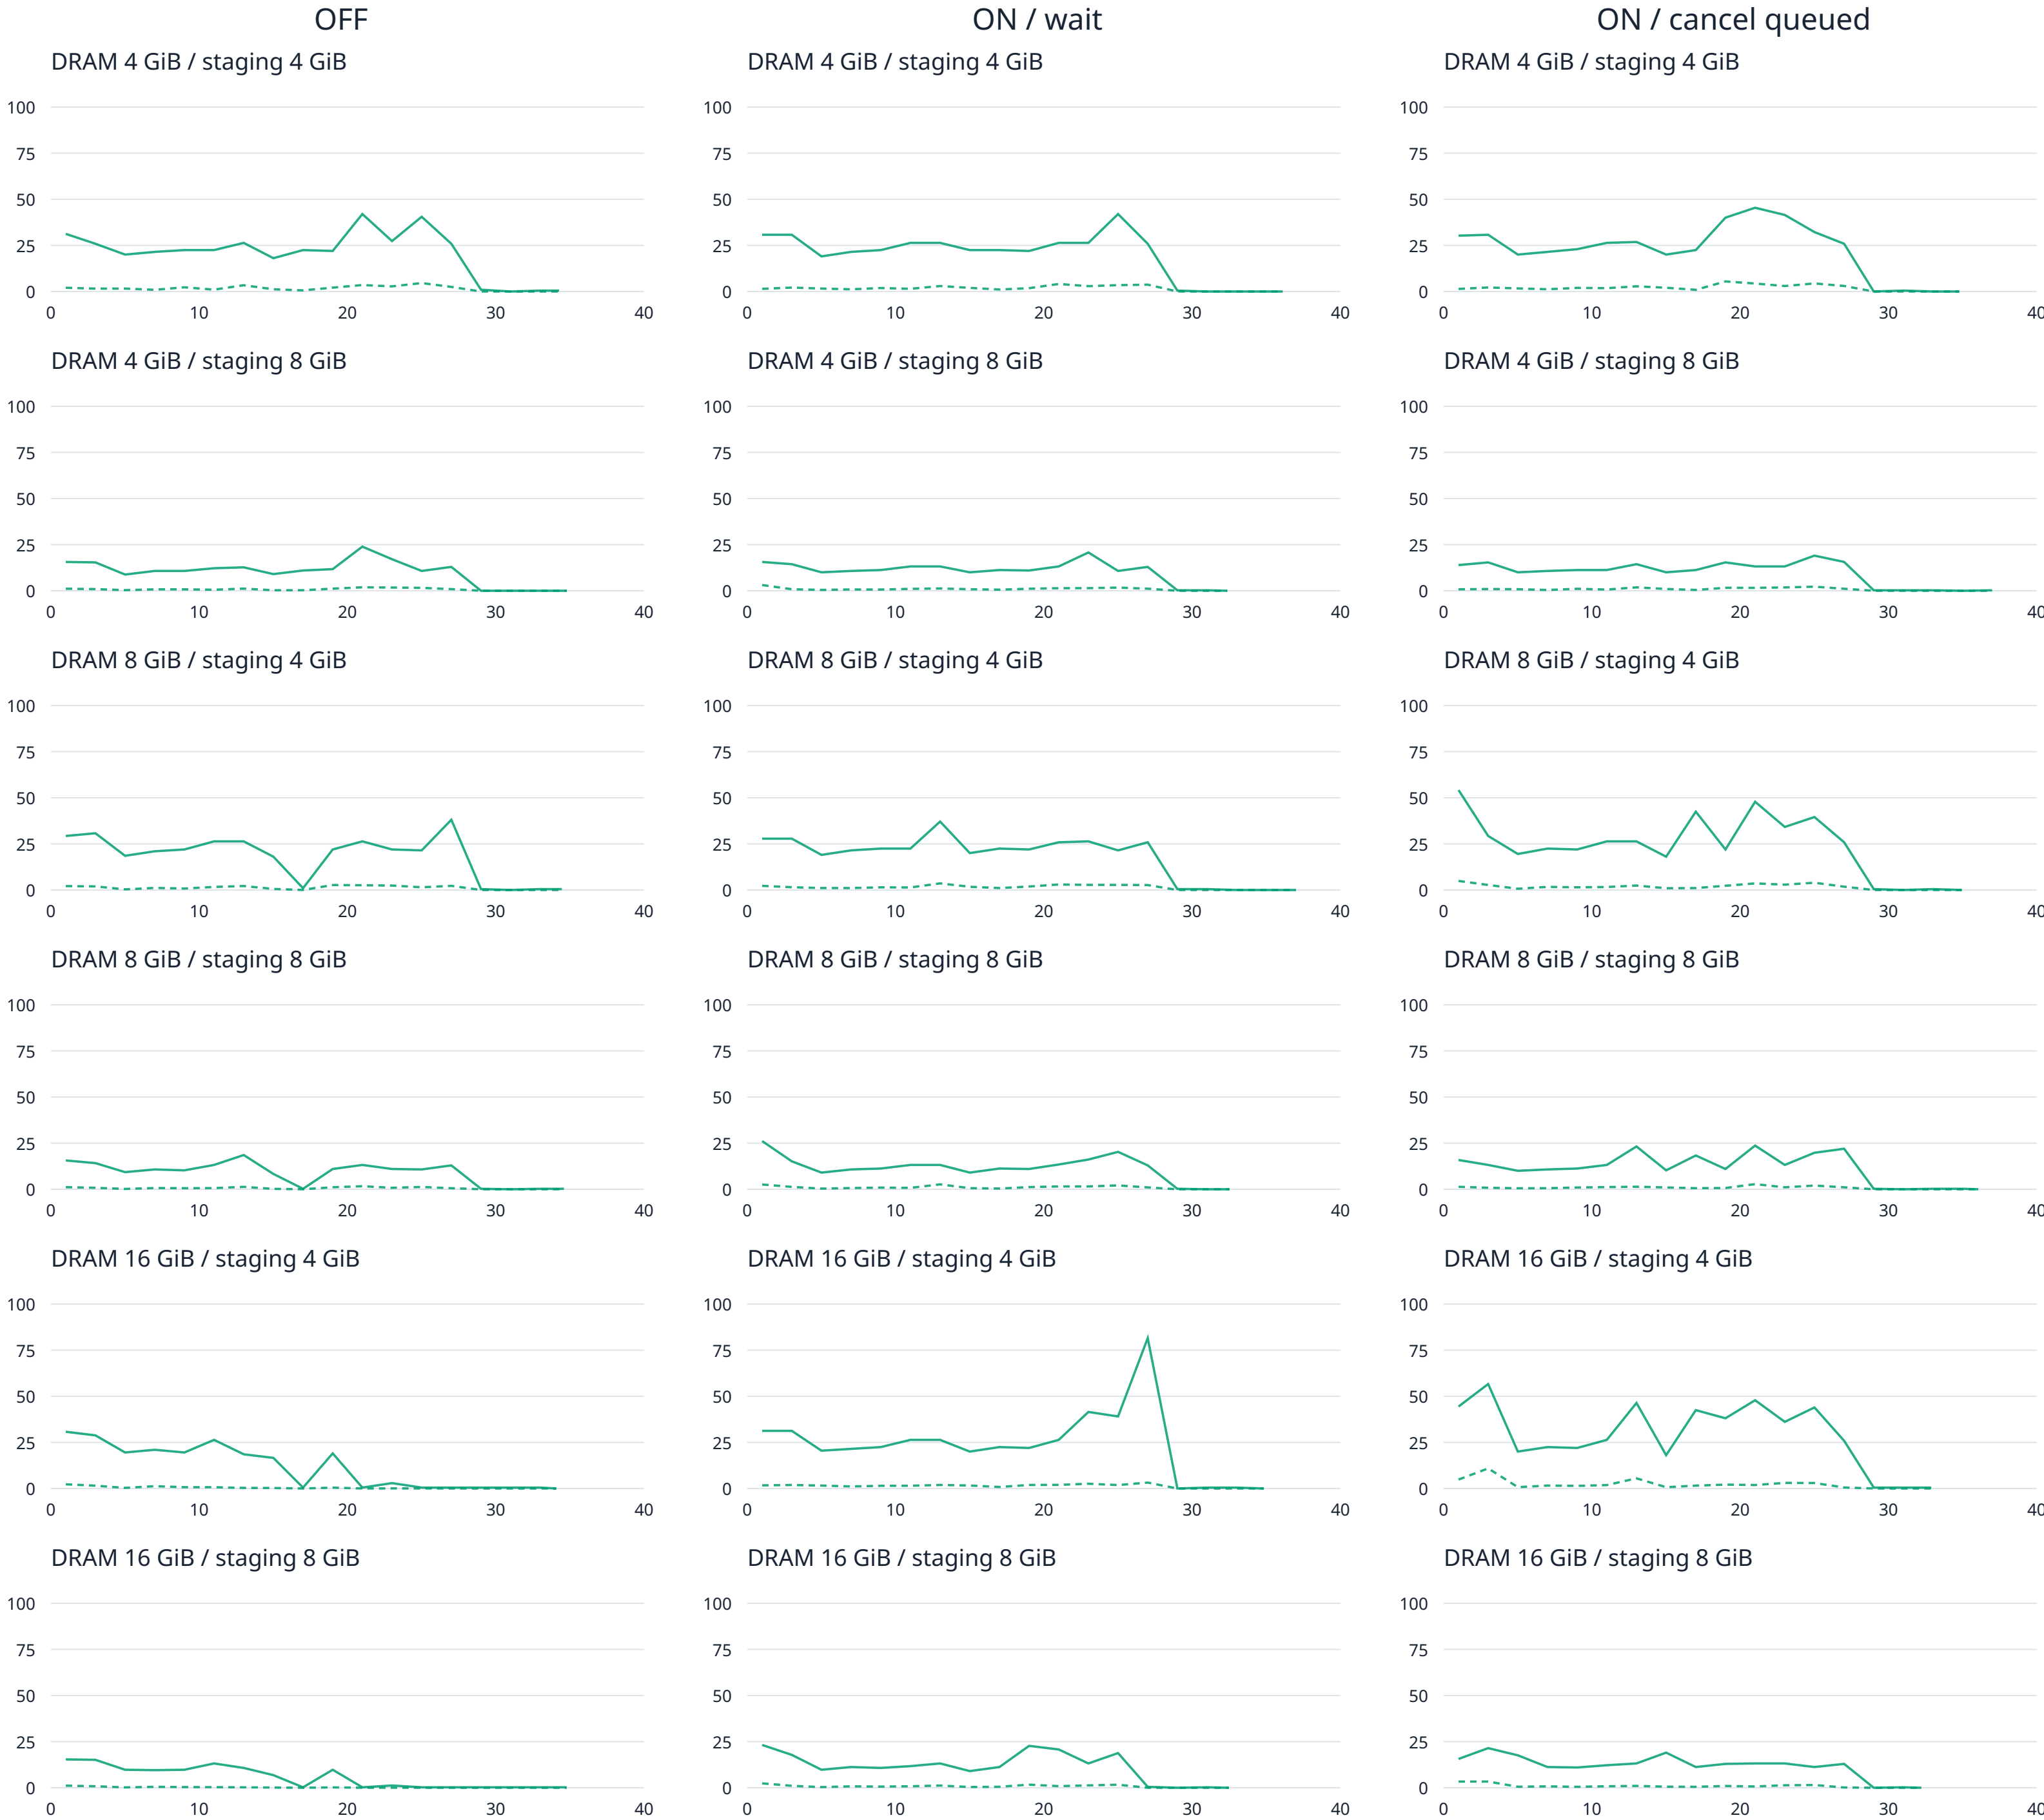


All panels share the same axes. Runs end at different times; no zero padding or time normalization.
Sampled peaks may miss brief allocation spikes. Occupancy is not GPU compute utilization.
OFF disables only DRAM prefetch; DAOS prefetch remains enabled. Equal times do not imply equal requests.

L-Eval | GPU staging occupancy | concurrency 8 | repeat 3

Qwen3-14B | 72 requests/run | 2 DRAM prefetch workers | 18 documents x 4 questions

Repeat 1 Repeat 2 Repeat 3 Solid: sampled peak per 2s bin | Dashed: sampled mean per 2s bin
Y: allocated staging / configured capacity (%) X: seconds since each run started

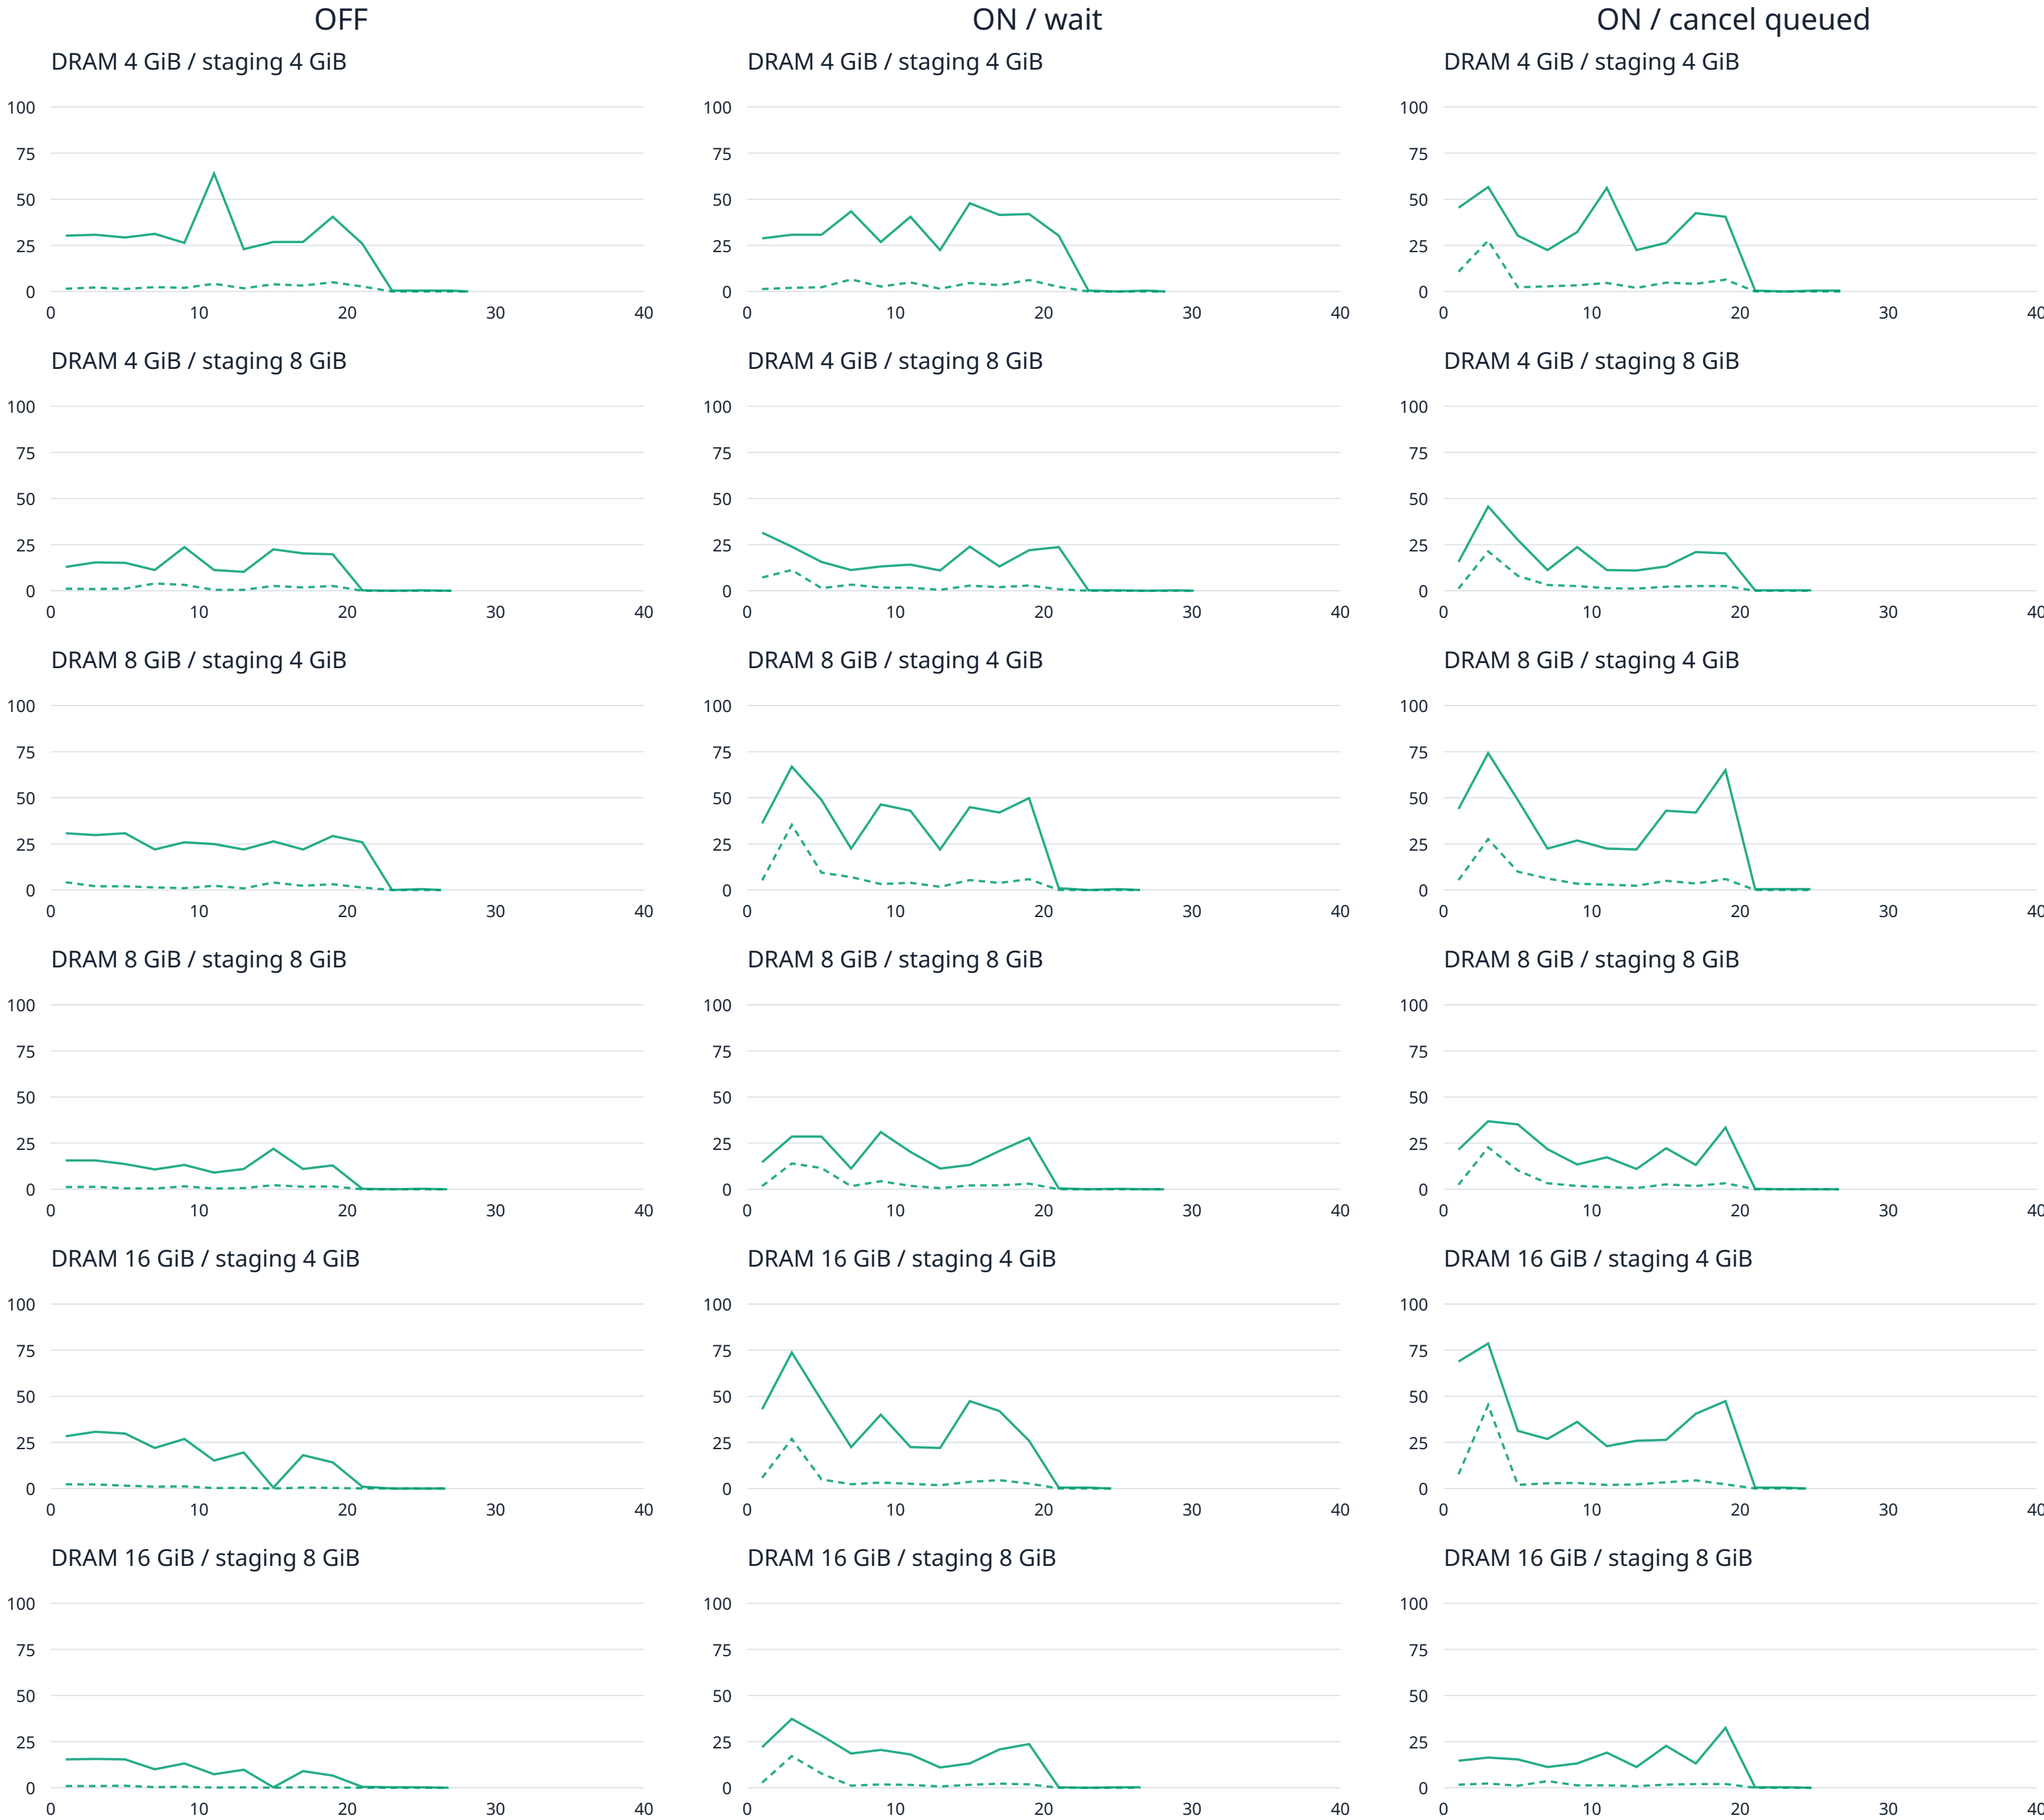


All panels share the same axes. Runs end at different times; no zero padding or time normalization.
Sampled peaks may miss brief allocation spikes. Occupancy is not GPU compute utilization.
OFF disables only DRAM prefetch; DAOS prefetch remains enabled. Equal times do not imply equal requests.

L-Eval | GPU staging occupancy | concurrency 16 | repeat 3

Qwen3-14B | 72 requests/run | 2 DRAM prefetch workers | 18 documents x 4 questions

Repeat 1 Repeat 2 Repeat 3 Solid: sampled peak per 2s bin | Dashed: sampled mean per 2s bin
Y: allocated staging / configured capacity (%) X: seconds since each run started



All panels share the same axes. Runs end at different times; no zero padding or time normalization.
Sampled peaks may miss brief allocation spikes. Occupancy is not GPU compute utilization.
OFF disables only DRAM prefetch; DAOS prefetch remains enabled. Equal times do not imply equal requests.